

Dynamic Web Development

Produced
by

David Drohan (ddrohan@wit.ie)

Department of Computing & Mathematics
Waterford Institute of Technology

<http://www.wit.ie>



Waterford Institute of Technology
INSTITIÚID TEICNEOLAÍOCHTA PHORT LÁIRGE





SERVER SIDE JAVASCRIPT

Outline

1. Introduction – What Node is all about
2. Events – Nodes Event-Driven, Non-Blocking I/O model
3. Node Modules – The Building Blocks of Node
4. Express – A Framework for Node
5. REST – The Architectural Style of the Web
6. API Design – Exposing Application Functionality
7. REST in Express – Leveraging URLs, URI's and HTTP
8. Demo – Labs in action

Outline

1. Introduction – What Node is all about
2. Events – Nodes Event-Driven, Non-Blocking I/O model
3. Node Modules – The Building Blocks of Node
4. Express – A Framework for Node
5. REST – The Architectural Style of the Web
6. API Design – Exposing Application Functionality
7. REST in Express – Leveraging URLs, URI's and HTTP
8. Demo – Labs in action

Introduction

WHAT NODE IS ALL ABOUT

Background

V8 is an open source JavaScript engine developed by Google. Its written in C/C++ and is used in Google Chrome Browser (and is fast!)

Node.js runs on V8.

It was created by **Ryan Dahl** in 2009.

Finally out of Beta phase.

- Latest LTS Version is v4.2.6
- Latest stable version is v5.5.0 (nodejs.org) (as at 27/01/16)

Is **Open Source**. It runs well on Linux systems, can also run on Windows systems.

It comes with a built-in **HTTP server library**

It has lots of libraries and tools available; even has its own **package manager (npm)**

Introduction: Basic

In simple words Node.js is 'server-side JavaScript'.

In not-so-simple words Node.js is a **high-performance** network applications framework, well optimized for high concurrent environments.

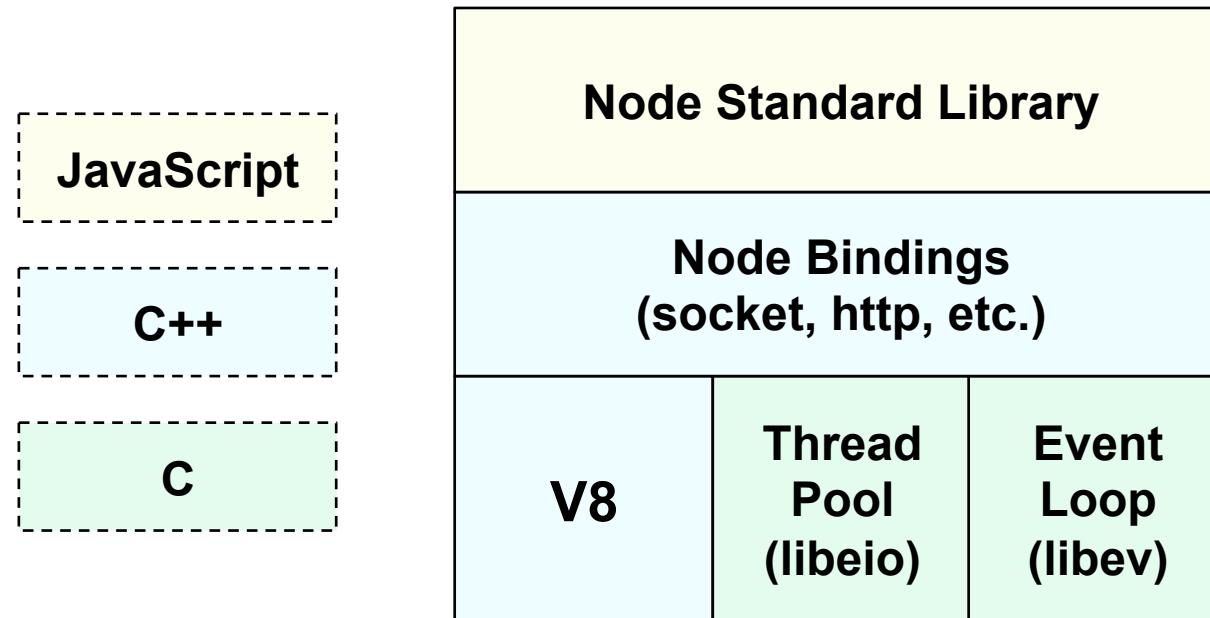
It's a **command line** tool.

In 'Node.js', '.js' doesn't mean that its solely written in JavaScript. It is 40% JS and 60% C++. (next slide)

From the official site:

'Node's goal is to provide an easy way to build scalable network programs' - (from nodejs.org!)

Introduction: Node Architecture



Introduction: Advanced (& Confusing)

Node.js uses an **event-driven, non-blocking I/O** model, which makes it lightweight. (again, from nodejs.org!)

It makes use of **event-loops** via JavaScript's **callback** functionality to implement the **non-blocking I/O**.

Programs for Node.js are written in JavaScript but not in the same JavaScript we are use to. There is **no DOM implementation** provided by Node.js, i.e. you **can not** do this:

```
var element = document.getElementById("elementId");
```

Everything inside Node.js runs in a **single-thread** (which must **never block!**).

If your program needs to wait for something (e.g., a response from some server you contacted), it **must provide a callback function**

NODE IS DEPLOYED BY BIG BRANDS

Big brands are using Node to power their business

Manufacturing	Financial	eCommerce	Media	Technology
    SIEMENS	  Goldman Sachs PayPal 	amazon.com  ebay TARGET Zappos.com	 CONDÉ NAST DOW JONES The New York Times SONY	salesforce.com  box  YAHOO!

Local Company

And most recently....



When to use Node.js?

Node.js is good for creating streaming based real-time services, web chat applications, static file servers etc.


If you need high level concurrency and not worried about CPU-cycles.

If you are great at writing JavaScript code because then you can use the same language at both the places: *server-side and client-side*.

More can be found at:

<http://stackoverflow.com/questions/5062614/how-to-decide-when-to-use-nodejs>

Some Node.js benchmarks




Taken from:

<http://code.google.com/p/node-js-vs-apache-php-benchmark/wiki/Tests>

A benchmark between Apache+PHP and node.js, shows the response time for 1000 concurrent connections making 10,000 requests each, for 5 tests.

Taken from: <http://nodejs.org/jsconf2010.pdf>
The benchmark shows the response time in milli-secs for 4 evented servers.



When to not use Node.js

When you are doing heavy and CPU intensive calculations on server side, because event-loops are CPU hungry.

Node.js API is finally out of beta, but it will keep on changing from one revision to another and there is a very little backward compatibility. A lot of the packages are also unstable. Therefore is not production ready just yet.

Node.js is a no match for enterprise level application frameworks like Spring(java), Django(python), Symfony/php) etc. Applications written on such platforms are meant to be highly user interactive and involve complex business logic.

Read further on disadvantages of Node.js on Quora:

<http://www.quora.com/What-are-the-disadvantages-of-using-Node-js>

Events


NODES EVENT-DRIVEN, NON-BLOCKING I/O MODEL

Some Theory: Events

- Generally, input/output (I/O) is slow.
 - Reading/writing to data store, probably across a network.
- Calculations in cpu are fast.
 - $2+2=4$
- Most time in programs is spent waiting for I/O to complete.
 - In applications with lots of concurrent users (e.g. web servers), you can't just stop everything and wait for I/O to complete.
- Solutions to deal with this are:
 - Blocking code with multiple threads of execution (e.g. Apache, IIS Servers)
 - Non-blocking, event-based code in single thread (e.g. NGINX, Node.js Servers)

Some Theory: Event-loops

Event-loops are the core of event-driven programming, almost all the UI programs use event-loops to track the user event, for example: Clicks, Ajax Requests etc.



Node Standard Library		
Node Bindings (socket, http, etc.)		
V8	Thread Pool (libeio)	Event Loop (libev)

Some Theory: Event-loops



Warning! Be careful to keep CPU intensive operations off the event loop.

Some Theory: Non-Blocking I/O

Traditional I/O


```
var result = db.query("select x from table_Y");  
doSomethingWith(result); //wait for result!  
doSomethingWithoutResult(); //execution is blocked!
```

Non-traditional, Non-blocking I/O

```
db.query("select x from table_Y",function (result){  
  doSomethingWith(result); //wait for result!  
});  
doSomethingWithoutResult(); //executes without any delay!
```


Blocking (Traditional)

- Traditional code waits for input before proceeding (Synchronous)
- The thread on a server "blocks" on I/O and resumes when it returns.



Non-blocking (Node)

- Node.js code runs in a Non-blocking (Asynchronous), event-based Javascript thread
 - No overhead associated with threads
 - Good for high concurrency (i.e. lots of client requests at the same time)



Blocking vs. Non-blocking

- Threads consume resources
 - Memory on stack
 - Processing time for context switching etc.
- No thread management on single threaded apps
 - Just execute “callbacks” when event occurs
 - Callbacks are usually in the form of anonymous functions.

Blocking | I/O Model

Example: ways in which a server can process orders from customers

*Hi, my name is Apache.
How may I take your
order?*



- The server serves one customer at a time.
- As each customer is deciding on their order, the server **sits and waits**.
- When the customer decides on an order, the server processes their order and moves on to the next customer.

Blocking | I/O Model



Blocking | I/O Model



Pseudocode:

```
order1 = db.query("SELECT * FROM  
menu WHERE preference = most")
```

```
order1.process
```

```
order2.process
```


Blocking | I/O Model



The more customers you want to serve at once, the more cashier lines you'll need.

Cashier lines ~ threads in computing

Multi-threaded processing

Parallel code execution

Multiple CPUs run at a time, utilizing shared resources (memory)

Non-Blocking | I/O Model



- Node loops through the customers and polls them to determine which ones are ready to order.
- During a function's queue, Node can listen to another event.
- When the other customer is finally ready to order, he'll issue a *callback*.
- Asynchronous **callbacks**: “*come back to me when I'm finished*”
 - function called at the completion of a given task.

Non-Blocking | I/O Model



Node code

```
console.log('Hello');

setTimeout(function () {
  console.log('World');
}, 5000);

console.log('Bye');
```

// Outputs:
// Hello
// Bye
// World

Allows for high concurrency

Non-Blocking | I/O Model



Every function in Node is non-blocking

Single-threaded

No parallel code execution


Single CPU

Non-Blocking | I/O Model

concurrency benchmark

concurrency	100	200	400	800
Apache PHP	6.337	6.955	6.723	19.232
nodeJS	3.461	3.284	3.721	3.689

Node is great for applications with high concurrency




(Concurrency = number of concurrent clients or users)

Non-Blocking | I/O Model

nginx: non-blocking I/O

apache: blocking I/O



Callbacks

In a *synchronous* program, you would write something along the lines of:

```
function processData () {  
  var data = fetchData ();  
  data += 1;  
  return data;  
}
```

This works just fine and is very typical in other development environments.

However, if *fetchData* takes a long time to load the data, then this causes the whole program to 'block' - otherwise known as sitting still and waiting - until it loads the data.

Node.js, being an *asynchronous* platform, doesn't wait around for things like file I/O to finish - Node.js uses **callbacks**.

Callbacks

If Google's V8 Engine is the heart of your Node.js application, then **callbacks** are its veins.

They enable a balanced, non-blocking flow of asynchronous control across modules and applications.

But for callbacks to work at scale you need a common, reliable protocol.

The “***error-first***” callback (also known as an “errorback”, “errback”, or “node-style callback”) was introduced to solve this problem, and has since become the standard for Node.js callbacks.

A **callback** is basically a function called at the completion of a given task; this prevents any blocking, and allows other code to be run in the meantime.

Defining an Error-First Callback

There's really only two rules for defining an error-first callback:

The first argument of the callback is reserved for an error object. If an error occurred, it will be returned by the first err argument.

The second argument of the callback is reserved for any successful response data. If no error occurred, err will be set to null and any successful data will be returned in the second argument.

```
fs.readFile('/foo.txt', function(err, data) {  
  // TODO: Error Handling Still Needed!  
  console.log(data);  
});
```

Defining an Error-First Callback

`fs.readFile()` takes in a file path to read from, and calls your **callback** once it has finished.

If all goes well, the file contents are returned in the **data** argument.

But if something goes wrong (the file doesn't exist, permission is denied, etc) the first **err** argument will be populated with an error object containing information about the problem.

It's up to you, the callback creator, to properly handle this error. You can throw an error if you want your entire application to shutdown. Or if you're in the middle of some asynchronous flow you can propagate that error out to the next callback. The choice depends on both the situation and the desired behavior.

Defining an Error-First Callback

```
fs.readFile('/foo.txt', function(err, data) {  
  // If an error occurred, handle it (throw, propagate, etc)  
  if(err) {  
    console.log('Unknown Error');  
    return;  
  }  
  // Otherwise, log the file contents  
  console.log(data);  
});
```

Callbacks

The node.js way to deal with the previous example we saw would look a bit more like this:

```
function processData (callback) {  
  fetchData(function (err, data) {  
    if (err) {  
      console.log("An error has occurred. Abort everything!");  
      callback(err);  
    }  
    data += 1;  
    callback(data);  
  });  
}
```

Callbacks

At first glance, it may look unnecessarily complicated, but **callbacks** are *the foundation of Node.js*.

Callbacks give you an interface with which to say, "and when you're done doing that, do all this." This allows you to have as many IO operations as your OS can handle happening at the same time.

For example, in a web server with hundreds or thousands of pending requests with multiple blocking queries, performing the blocking queries asynchronously gives you the ability to be able to continue working and not just sit still and wait until the blocking operations come back.

This is a major improvement.

Node Modules

THE BUILDING BLOCKS OF NODE

Node.js Ecosystem

Node.js relies heavily on **modules**.

Creating a module is easy, just put your JavaScript code in a separate js file and include it in your code by using the keyword **require**, like:

```
var modulex = require('./modulex');
```

Libraries in Node.js are called packages and they can be installed by typing

```
npm install "package_name"; //installs in current folder  
//package should be available in npm registry @ nmpjs.org
```

NPM downloads and installs modules, placing them into a **node_modules** folder in your current folder.



NPM

- Common npm commands:
 - **npm init** initialize a package.json file
 - **npm install <package name> -g** install a package, if **-g** option is given package will be installed globally, **--save** and **--save-dev** will add package to your dependencies
 - **npm install** install packages listed in package.json
 - **npm ls -g** listed local packages (without **-g**) or global packages (with **-g**)
 - **npm update <package name>** update a package



Creating your own Node Modules donations.js

```
var donations = require('../models/donations');
var express = require('express');
var router = express.Router();
```

```
function getByValue(arr, id) {...}
```

```
router.findAll = function(req, res) {
  // Return a JSON representation of our list
  res.json(donations);
}
```

```
router.findOne = function(req, res) {...}
```

```
router.addDonation = function(req, res) {...}
```

```
router.deleteDonation = function(req, res) {...}
```

```
router.incrementUpvotes = function(req, res) {...}
```

```
module.exports = router;
```

```
var routes = require('./routes/index');
var donations = require('./routes/donations');
var app = express();
```

app.js

```
//Our Custom Routes
app.get('/donations', donations.findAll);
app.get('/donations/:id', donations.findOne);
app.post('/donations', donations.addDonation);
app.put('/donations/:id/votes', donations.incrementUpvotes);
app.delete('/donations/:id', donations.deleteDonation);
```

app.js

Defines what
'require' returns

The require search

- Require searches for modules based on path specified:
- Just providing the module name will search in node_modules folder

```
var myMod = require('./myModule'); //current dir  
var myMod = require('../myModule'); //parent dir  
var myMod = require('../../modules/myModule');
```

```
var myMod = require('myModule');
```

Express

A FRAMEWORK FOR NODE

What is Express?

Express is a minimal and flexible framework for writing web applications in Node.js

- Built-in handling of HTTP requests
- You can tell it to 'route' requests for certain URLs to a function you specify
 - Example: When /login is requested, call function handleLogin()
- These functions are given objects that represent the request and the response, not unlike Servlets
- Supports parameter handling, sessions, cookies, JSON parsing, and many other features
- API reference: <http://expressjs.com/api.html>

```
var express = require('express');
var app = express();

app.get('/', function(req, res) {
  res.send('hello world');
});

module.exports = app;
```

We'll cover this (and more) in detail in the next Section (Part 2)

Using NodeJS

NodeJS is *just* a JavaScript interpreter.

It comes with a package manager called npm

- Install packages like this: `npm install <package_name>`
- This will install it in the current folder.
- To install globally, do `npm install -g <package_name>`

To use Node as a webserver, you must write an application that responds to web requests.

Node has a library (HTTP) for doing this, but it's easier to use a framework, like *Express*

To access a library, use the `require()` function

Using Express

Express is just a package for Node

- Create a new web application with `var app = express();`
- Respond to requests like `app.get('/user', function(req, res) {`
- Look at parameters through the `req` object
 - `req.params` for query parameters
 - `req.body` for post fields
 - `req.files` for files
- Send responses through the `res` object
 - `res.send("Hi mom!")`
- Start the application with
 - `app.listen(<port>)`

Great Resources

Official Tutorial – <https://nodejs.org/documentation/tutorials/>

Official API – <https://nodejs.org/api/>

Developer Guide – <https://nodejs.org/documentation>

Video Tutorials – <http://nodetuts.com>

Video Introduction – <https://www.youtube.com/watch?v=FqMlyTH9wSg>

YouTube Channel – https://www.youtube.com/channel/UCvhlsEIBlfWSn_Fod8FuuGg

Articles, explanations, tutorials – <https://nodejs.org/community/>

Questions?